

Arsenii Ashukha

[Home page](#) / [Google Scholar](#) / [GitHub](#)

e-mail: ars.ashuha@gmail.com tel: +79651782915

I'm a Research Scientist at Samsung AI Center Moscow. I (almost) received a PhD in Bayesian deep learning [@bayesgroup](#) so I can make big overcomplicated DNNs work :) My research interests are focused on crafting and training reliable and data&compute-efficient deep neural networks.

My works are the best known for:

- **Sparse variational dropout** (ICML'17)—a method for sparsification of DNNs that archives an extreme level of CNN pruning (up to 250x) with no accuracy loss, SparseVD-like approaches were used in many other sparsification algorithms, and models beyond CNNs e.g., [Transformers](#).
- **Pitfalls of in-domain uncertainty estimation** work (ICLR'20), that demonstrates that Deep Ensemble—a pioneering method of the field—is still the best solution for in-domain uncertainty estimation, and many recent (Bayesian) techniques fall behind in pure performance and performance per training or inference compute. While simple test-time augmentation allows significantly improve Deep Ensembles for free.

Prior to Samsung, I was a part of Yandex Research (a Russian search giant) in collaboration with University of Amsterdam (with Dmitry Vetrov and Max Welling) and did ML engineering internships at Yandex and Rambler.

SELECTED PUBLICATIONS

Google Scholar: scholar.google.com/citations?user=IU-kuP8AAAAJ

Most representative works are **highlighted**

- ***Pitfalls of In-Domain Uncertainty Estimation and Ensembling in Deep Learning***, ICLR 2020 [[arXiv](#)]
[Arsenii Ashukha](#), Alexander Lyzhov, Dmitry Molchanov, Dmitry Vetrov
- ***Variational Dropout Sparsifies Deep Neural Networks***, ICML 2017 [[arXiv](#)]
Dmitry Molchanov, [Arsenii Ashukha](#), Dmitry Vetrov
- ***Greedy Policy Search: A Simple Baseline for Learnable Test-Time Augmentation***, UAI 2020 [[arXiv](#)]
Dmitry Molchanov, Alexander Lyzhov, Yuliya Molchanova, [Arsenii Ashukha](#), Dmitry Vetrov
- ***The Deep Weight Prior***, ICLR 2019 [[arXiv](#)]
Andrei Atanov, [Arsenii Ashukha](#), Kirill Struminsky, Dmitry Vetrov, Max Welling
- ***Variance Networks: When Expectation Does Not Meet Your Expectations***, ICLR 2019 [[arXiv](#)]
Kirill Neklyudov, Dmitry Molchanov, [Arsenii Ashukha](#), Dmitry Vetrov
- ***Uncertainty Estimation via Stochastic Batch Normalization***, ICLR Workshop 2018 [[arXiv](#)]
Andrei Atanov, [Arsenii Ashukha](#), Dmitry Molchanov, Kirill Neklyudov, Dmitry Vetrov
- ***Structured Bayesian Pruning via Log-Normal Multiplicative Noise***, NeurIPS 2017 [[arXiv](#)]
Kirill Neklyudov, Dmitry Molchanov, [Arsenii Ashukha](#), Dmitry Vetrov

CODE

- Check out very short and simple and fun to make implementations of ML algorithms:
 - [Gradient Boosting](#)
 - [Real NVP normalizing flows](#)
 - [Quantile Regression DQN \(Distributional RL\)](#)
- Also, check out more solid implementations:
 - [Multi-GPU SimCLRv1](#) closely reproduced results on both CIFAR-10 and ImageNet
 - [Ensembles \(Deep ensembles, Snapshot ensembles, cSGLD, FGE, etc.\)](#)

EDUCATION

- 2017-2021 PhD in Computer Science, **Centre of Deep Learning National Research University HSE**. My PhD is focused on *applications and understanding of stochastic deep learning models*, Advisor: [Dmitry Vetrov](#).
- 2015-2017 MSc in Computer Science, **Moscow Institute of Physics and Technology** jointly with **Yandex School of Data Analysis**, where I worked with [Dmitry Vetrov](#) on learning sparse DNNs.
- 2011-2015 BSc in Computer Science **Bauman Moscow State Technical University** with a double major in applied math and CS. I worked on linguistically motivated topic models with [Natalia Loukachevitch](#).

PROFESSIONAL EXPERIENCE

- *Research Scientist, Samsung AI Center* (2018 - Now):

Showed that modern uncertainty estimation techniques perform worse than the simplest 5 year old baseline, and widely used metrics for asses uncertainty estimation cannot provide a fair ranking of methods (results published at ICLR'20).

The other topics I worked/working on include learnable data augmentation, learning prior distribution over DNNs' kernels with generative models, and self/semi-supervised learning. Since 2021 I'm helping to manage the lab.

- *Research Scientist, Yandex.Research & University of Amsterdam* (2016 - 2018):

Created Sparse Variational Dropout a method for sparsification of DNNs that, for the first time, allowed to achieve over 250x compression ratio (results published at ICML'17). The modified version of SparseVD with neuron-level sparsity allowed to accelerate inference of a CNN by 2-5 times and was applied to the feature extraction process for image retrieval (results published at NeurIPS).

- *Machine Learning Engineer Intern, Yandex.Music* (summer of 2016):

I worked on feature extraction techniques for music data with convolutional neural networks. I also developed an evaluation of learned representations. The representations were used in the content-based recommendation system for yandex music.

- *Machine Learning Engineer Intern, Rambler&Co* (May - Oct 2015):

Worked on demographic classification and recommendation systems. My responsibility included improving the quality and performance of classifiers, automatic feature extraction algorithms, and recommendation algorithms. Stack of technologies: Hadoop, Hive, Spark, XGboost, VW, gensim.

REVIEWING

- Conferences:
 - International Conference on Machine Learning, ICML (2019, 2020 top-33% highest-scored reviewers)
 - Neural Information Processing Systems, NeurIPS 2019 (top-50% highest-scored reviewers)
 - International Conference on Learning Representations, ICLR (2020, 2021)
- Workshops:
 - ICML Workshop on Invertible Neural Networks (2019, 2021, invertibleworkshop.github.io)
 - Bayesian Deep Learning Workshop (since 2017, bayesiandeeplearning.org)

TEACHING

- Supervisor of reading clubs on machine learning at HSE and Yandex school of data analysis (since 2017)
- Talks and practical sessions at **Deep | Bayes** Summer School on Bayesian Deep Learning (since 2017)
- Lecturer, Moscow Institute of Physics and Technology: I was a lecturer and manager of the deep learning brunch of a faculty-wide machine learning course ~60 students (ml-mipt.github.io). Also, I taught deep learning and practical sessions on cutting-edge ML algorithms on a facultative course "Data Mining in Action" ~ 200 students (<https://bit.ly/3eRLGYp>). The goal of this course is to make ML education available for everyone for free.

SUPERVISION

- Alexander Lyzhov (moved to The Center on Long-Term Risk), Deep Neural Network Ensembles: Analysis and Approaches to Diversification (MSc, 2020)
- Andrei Atanov (moved to EPFL, PhD student), Effective Learning of Deep Neural Networks Ensembles (BSc, 2018), Learning Deep Models with Small Data (MSc, 2020)
- Evgenii Nikishin (moved to Mila, PhD student), Stability Improvement and Knowledge Transfer in Deep Reinforcement Learning (MSc, 2019)

FRAMEWORKS & PROGRAMMING LANGUAGES

- I'm fluent in Python and I use to code in C/C++, Go, language is not a problem after all.
- I'm also fluent with common Python libs such as NumPy, Matplotlib, scikit-learn, etc.
- My primary deep learning framework at the moment is PyTorch. Prior to that, I had a decent experience with Theano+Lagange and TensorFlow.
- I have experience with MapReduce, Hadoop, Hive, and Spark.