

Scientific Seminar "Bayesian Methods of ML"

Deep Structured Models

Ashuha Arseniy


Moscow Institute of Physics and Technology

`ars.ashuha@gmail.com`

April 28, 2016

Based on article: *Chen, Schwing, et al. "Learning deep structured models." ICML 2015*

- 1 Neural Nets and Graphical Models
- 2 Deep Structured Models
- 3 Efficient Approximate Learning of DSM
- 4 Blending Learning
- 5 Evaluation



- ▶ **NNs** is a framework for constructing flexible models
- ▶ Neural net is a **composition** linear and nonlinear functions


$$\operatorname{argmax}(\sigma[\operatorname{Linear}(\sigma[\operatorname{Linear}(\text{img}, w)], w)]) = \text{cat}$$

- ▶ **We can learn it efficiently** by back propagation

Problem

Can't take into account dependences between predicted variables.

- ▶ Non structured – predict simple variable (like a number)
- ▶ Structured – predict difficult variable (like a matrix, tree, sequence)




(a) Segmentation, $|Y| = \#pixel \times \#segment$




water/animals/sea

(b) Tagging, $|Y| = 2^{\#tags}$




(c) Traffic prediction, $|Y| = ?$



(d) Translation, $|Y| = ?$

- ▶ **GMs** is a framework for taking into account dependencies between predicted variables.
- ▶ What can we do with exp-large output space? Use local dependences.
- ▶ Introduce prior knowledge as a score functions $\phi_r(y_r)$, $|y_r|$ is small



$\phi_{y1,y2}(\text{people}, \text{male}) = 10$ is high
 $\phi_{y1,y2}(\text{wather}, \text{girl}) = 0.2$ is low

- ▶ We can introduce non-normalized probability distribution over outputs

$$p(y|x, w) = \frac{1}{Z} \prod_r \phi_r(x, y_r; w) \quad \text{Energy} = - \sum_r \phi_r(x, y_r; w)$$

- ▶ **Inference Task:**

$$y^* = \arg \max_y p(y|x, w)$$

1. We want to **train parameters** w of parametric potential
2. Given training data $(x, y) \in D$; estimate the functions $f_r(y, x, w)$
3. Minimize a typically convex loss and a regularize on training set

$$Loss_{log}(x, y, w) = -\ln p_{x,y}(y; w)$$

$$Loss_{hinge}(x, y, w) = \max_{\hat{y}} (\Delta(y, \hat{y}) - w^T \Phi(x, \hat{y}) + w^T \Phi(x, y))$$

4. The assumption is that the model is **log-linear**

$$E(x, y, w) = -w^T \phi(x, y)$$

and the features decompose in a graph

$$w^T \phi(x, y) = \sum_{r \in R} w_r^t \phi_r(x, y)$$

Problem

How can we remove the log-linear restriction,
to use potentials such as Neural Nets?

- 1 Neural Nets and Graphical Models
- 2 Deep Structured Models**
- 3 Efficient Approximate Learning of DSM
- 4 Blending Learning
- 5 Evaluation


How can we combine Graphical Models and Deep Neural Nets?

1. Peace-wise learning:

- ▶ train deep features \rightarrow train linear potential \rightarrow inference in GM

2. Jointly learning:

- ▶ train deep features as non linear potential \rightarrow inference in GM




How can we combine Graphical Models and Deep Neural Nets?

1. Peace-wise learning:

- ▶ train deep features \rightarrow train linear potential \rightarrow inference in GM

2. Jointly learning:

- ▶ train deep features as non linear potential \rightarrow inference in GM



- ▶ We have: scoring function $F(x, y; w)$, training data $(x, y) \in D$
- ▶ **Prediction proses** is equal finding maximum scoring configuration y^* :

$$y^* = \arg \max_y F(x, y; w)$$

- ▶ Introduce probability distribution over configurations as

$$p_{(x,y)}(\hat{y}|w) = \frac{\exp F(x, \hat{y}, w)}{\sum_{y'} \exp F(x, y', w)} = \frac{\exp F(x, \hat{y}, w)}{Z(x, w)}$$

rephrase previous task as finding high probably configuration

- ▶ **Training proses** is finding parameters w by MLE

$$\begin{aligned} w &= \arg \max_w \log \prod_{(x,y) \in D} p_{(x,y)}(y|w) = \\ &= \arg \max_w \sum_{(x,y) \in D} F(x, y, w) - \ln Z(x, w) \end{aligned}$$

We have optimization problem:

$$\sum_{(x,y) \in D} \left(F(x, y, w) - \log \sum_{y' \in Y} \exp F(x, y', w) \right) \rightarrow \max_w$$

Let's solve it by gradient ascent (will be proof on the board if it's necessary):

$$\begin{aligned} & \frac{\partial}{\partial w} \sum_{(x,y) \in D} (F(x, y, w) - \log Z(x, w)) = \\ &= \sum_{(x,y) \in D} \sum_{y' \in Y} (p(y'|w, x) - \delta(y' = y)) \frac{\partial}{\partial w} F(x, y', w) \end{aligned}$$

Very easy! Where is a challenge?

Problem: What If Y is exponentially large!

1) How can we represent F ? 2) What we can do with sum over Y ?

$$\sum_{(x,y) \in D} \left(F(x, y, w) - \log \sum_{y' \in Y} \exp F(x, y', w) \right) \rightarrow \max_w$$

1. Use the graphical model $F(x, y; w) = \sum_r f_r(x, y; w)$

$$\begin{aligned} & \frac{\partial}{\partial w} \sum_{(x,y) \in D} (F(x, y, w) - \log Z(x, w)) = \\ &= \sum_{(x,y) \in D} \sum_{y' \in Y} (p(y'|w, x) - \delta(y' = y)) \frac{\partial}{\partial w} F(x, y', w) \end{aligned}$$

(will be proof on the board if it's necessary):

$$= \sum_{(x,y) \in D} \sum_{y'_r, r} (p_r(y'_r|w, x) - \delta(y'_r = y_r)) \frac{\partial}{\partial w} f_r(x, y'_r, w)$$

2. How to obtain marginals $p_r(y_r|w, x)$?
3. Use beliefs $p_r(y_r|w, x) \approx b_r(y_r|w, x)$

Deep Structured Learning (algo 1)

Repeat until stopping criteria:

1. Forward pass to compute the $f_r(y_r, x; w) \quad \forall r, y_r, (x, y) \in D$
2. Compute the $b_r(y_r|x, w)$ by approx inference $\forall r, y_r, (x, y) \in D$
3. Backward pass via chain rule to obtain gradient

$$\frac{\partial}{\partial w} = \sum_{(x, y) \in D, y'_r, r} (b_r(y'_r|w, x) - \delta(y'_r = y_r)) \frac{\partial}{\partial w} f_r(x, y'_r; w)$$

4. Update parameters w

$$w = w - \alpha \cdot \partial / \partial w$$

Problem

We run inference for each object to make one parameters update

- 1 Neural Nets and Graphical Models
- 2 Deep Structured Models
- 3 Efficient Approximate Learning of DSM**
- 4 Blending Learning
- 5 Evaluation

$$\sum_{(x,y) \in D} \left(F(x, y, w) - \log \sum_{y' \in Y} \exp F(x, y', w) \right) \rightarrow \max_w$$

1. We can represent Z as (will be proof on the board if it's necessary):

$$\ln Z = \sum_{\hat{y}} \exp F(x, \hat{y}, w) = \max_{p_{(x,y)}} \mathbb{E}_{p_{(x,y)}(\hat{y})} F(x, \hat{y}; w) + H(p_{(x,y)})$$

2. Assumption, F and H is decomposed into a sum of "local" functions

$$F = F(x, y; w) = \sum_r f_r(x, y_r; w) \quad H = H(p_{(x,y)}) = \sum_r H(p_{(x,y),r})$$

3. Rephrase our task as

$$\min_w \sum_{(x,y) \in D} \left(\max_{p_{(x,y)}} \left\{ \sum_r p_{(x,y),r}(\hat{y}_r) f_r(x, \hat{y}_r; w) + H(p_{(x,y)}) \right\} - F \right)$$

$$\sum_{(x,y) \in D} \left(\max_{p_{(x,y)}} \left\{ \sum_r p_{(x,y),r}(\hat{y}_r) f_r(x, \hat{y}_r; w) + H(p_{(x,y)}) \right\} - F \right) \rightarrow \min_w$$

1. We can't compute true marginals, let's use beliefs $b_{(x,y)} \approx p_{(x,y)}$

$$b_{(x,y)} \in C_{(x,y)} = \begin{cases} b_{(x,y),r}(\cdot) \geq 0 & \sum_{y_r} b_{(x,y),r}(y_r) = 1 & \forall r \\ b_{(x,y),r} = \sum_{\hat{y}_p \setminus \hat{y}_r} p_{(x,y),p}(\hat{y}_p) & & \forall r, \hat{y}_r, p \in P(r) \end{cases}$$

2. $P(r) = \{p \in Y : r \subset p\}$ and $C(r) = \{c \in Y : r \in P(c)\}$

3. Rephrase our task as

$$\min_w \sum_{(x,y) \in D} \left(\max_{b_{(x,y)} \in C_{(x,y)}} \left\{ \sum_r b_{(x,y),r}(\hat{y}_r) f_r(x, \hat{y}_r; w) + H(b_{(x,y)}) \right\} - F \right)$$

Problem

We need to solve inner problem to compute subgradient!

$$\min_w \sum_{(x,y) \in D} \left(\max_{b_{(x,y)} \in C_{(x,y)}} \left\{ \sum_r b_{(x,y),r}(y_r) f_r(x, y_r; w) + H(b_{(x,y)}) \right\} - F \right)$$

s.t. $b_{(x,y)} \in C_{(x,y)}$ marginalization and discrete distribution conditions

H is redefined as barrier function when argument is not a distribution:

1. The Lagrangian of inner problem is:

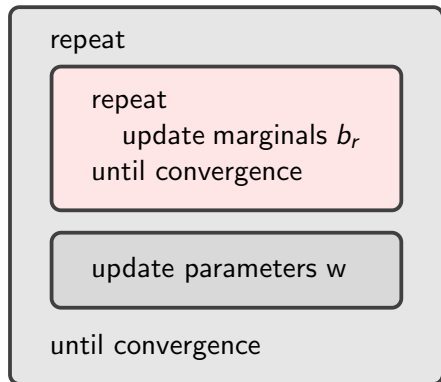
$$L_{(x,y)} = \sum_{r, \hat{y}_r} b_{(x,y),r}(\hat{y}_r) \cdot \hat{f}_r(x, \hat{y}_r; w, \lambda) + H_{barrier}$$

$$\hat{f}_r(x, \hat{y}_r; w, \lambda) = f_r(x, \hat{y}_r; w) + \sum_{p \in P(r)} \lambda_{(x,y),p \rightarrow r}(\hat{y}_r) - \sum_{c \in C(r)} \lambda_{(x,y),c \rightarrow r}(\hat{y}_c)$$

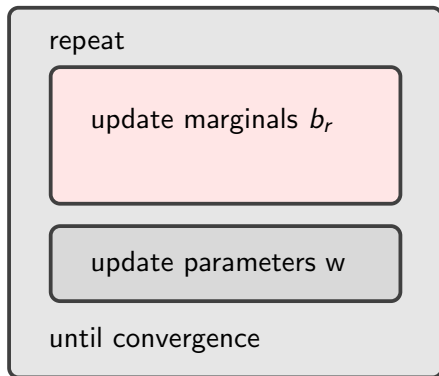
2. Move to dual task by λ ($\ln Z = \max_{p_{(x,y)}} \mathbb{E}_{p_{(x,y)}(\hat{y})} F(x, \hat{y}; w) + H(p_{(x,y)})$):

$$\min_{w, \lambda} \sum_{(x,y),r} \ln \sum_{\hat{y}_r} \exp \hat{f}_r(x, \hat{y}_r; w, \lambda) - \sum_{(x,y) \in D} F(x, y; w)$$

Standard learning:



Blended learning:



Advantage:

More frequent parameter updates

Hazan, Schwing, McAllester, Urtasun: Blending Learning and Inference in Structured Prediction

$$D(\lambda, w) = \sum_{(x,y),r} \ln \sum_{\hat{y}_r} \exp(f_r(x, \hat{y}_r; w) + \sum_{c \in C(r)} \lambda_{(x,y),c \rightarrow r}(\hat{y}_c) - \sum_{p \in P(r)} \lambda_{(x,y),r \rightarrow p}(\hat{y}_p)) - \sum_{(x,y)} F(x, y; w) \rightarrow \min_{w, \lambda}$$

- Optimize by w (will be proof on the board if it's necessary):

$$\frac{\partial D}{\partial w} = \sum_{(x,y),r,\hat{y}_r} b_{(x,y),r,\hat{y}_r} \frac{\partial}{\partial w} f_r(x, \hat{y}_r; w) + \sum_{(x,y)} \frac{\partial}{\partial w} F(x, y; w)$$

- Optimize by λ (will be proof on the board if it's necessary):

$$\mu_{(x,y),p \rightarrow r}(\hat{y}_r) = \ln \sum_{\hat{y}_p \in \hat{y}_r} \exp(f_p(x, \hat{y}_p; w) - \sum_{p' \in P(p)} \lambda_{(x,y),p \rightarrow p'}(\hat{y}_{p'}) + \sum_{r' \in C(p) \setminus r} \lambda_{(x,y),r' \rightarrow p}(\hat{y}_{r'}))$$

$$\lambda_{(x,y),r \rightarrow p}(\hat{y}_r) \propto c_r \cdot (f_r(x, \hat{y}_r; w) - \sum_{c \in C(r)} \lambda_{(x,y),c \rightarrow r}(\hat{y}_c) + \sum_{p \in P(r)} \mu_{(x,y),p \rightarrow r}(\hat{y}_r)) - \mu_{(x,y),p \rightarrow r}(\hat{y}_r)$$

Efficient Deep Structured Learning (algo 2)

Repeat until stopping criteria:

1. Forward pass to compute the $f_r(y_r, x; w) \quad \forall r, y_r, (x, y) \in D$
2. Compute the $b_r(y_r|x, w) = \exp(\hat{f}_r(x, y_r; w, \lambda)) \quad \forall r, y_r, (x, y) \in D, p \in P(r)$

$$\mu_{(x,y),p \rightarrow r}(\hat{y}_r) = \ln \sum_{\hat{y}_p \in \hat{P}(\hat{y}_r)} \exp(f_p(x, \hat{y}_p; w) - \sum_{p' \in C(p)} \lambda_{(x,y),p \rightarrow p'}(\hat{y}_p) + \sum_{r' \in C(p) \setminus r} \lambda_{(x,y),r' \rightarrow p}(\hat{y}_{r'}))$$


$$\lambda_{(x,y),r \rightarrow p}(\hat{y}_r) \propto c_r \cdot \left(f_r(x, \hat{y}_r; w) - \sum_{c \in C(r)} \lambda_{(x,y),c \rightarrow r}(\hat{y}_c) + \sum_{p \in P(r)} \mu_{(x,y),p \rightarrow r}(\hat{y}_r) \right) - \mu_{(x,y),p \rightarrow r}(\hat{y}_r)$$

3. Backward pass via chain rule to obtain gradient

$$g = \sum_{(x,y),r,\hat{y}_r} b_{(x,y),r}(\hat{y}_r) \nabla_w f_r(\hat{y}_r, x; w) - \nabla_w \sum_{(x,y),r} f_r(x, y; w)$$

4. Update parameters w

$$w = w - \alpha \cdot \partial / \partial w$$



1. Modeling of correlations between variables
2. Non-linear dependence on parameters
3. Joint training of many convolution neural networks

- 1 Neural Nets and Graphical Models
- 2 Deep Structured Models
- 3 Efficient Approximate Learning of DSM
- 4 Blending Learning
- 5 Evaluation**

- **Task:** Find a combination of tags that describe the image, $|Y| = 2^{38}$



female/indoor/portrait
female/indoor/portrait



sky/plant life/tree
sky/plant life/tree



water/animals/sea
water/animals/sky

- **Graphical Model:** Fully Connected 38
- **First order potential:** $f_i(x, y_i; U) = \text{Alexnet}(x, U)$
- **Second order potential:** $f_{i,j}(x, y_i, y_j; W) = W_{y_i y_j}$

Training method	Prediction error [%]
Unary only	9.36
Piecewise	7.70
Joint (with pre-training)	7.25

Learned class "correlations":

female	0.00	0.68	0.04	0.24	-0.01	-0.05	0.07	-0.01	0.01
people	0.68	0.00	0.06	0.36	-0.05	-0.12	0.74	-0.04	-0.03
indoor	0.04	0.06	0.00	0.07	-0.35	-0.34	0.02	-0.15	-0.21
portrait	0.24	0.36	0.07	0.00	-0.02	-0.01	0.12	0.02	0.05
sky	-0.01	-0.05	-0.35	-0.02	0.00	0.24	-0.00	0.44	0.30
plant life	-0.05	-0.12	-0.34	-0.01	0.24	0.00	-0.07	0.09	0.68
male	0.07	0.74	0.02	0.12	-0.00	-0.07	0.00	0.00	-0.02
clouds	-0.01	-0.04	-0.15	0.02	0.44	0.09	0.00	0.00	0.11
tree	0.01	-0.03	-0.21	0.05	0.30	0.68	-0.02	0.11	0.00

- **Task:** Find five letters within distorted images, $|Y| = 26^5$



banal



julep




resty

- **Graphical Model:**



1st order Markov





2nd order Markov

- **First order potential:**

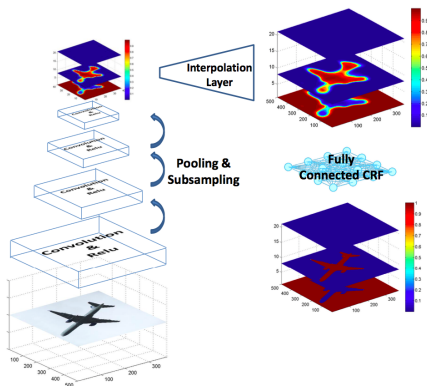
1. One Layer : $f_i(x, y_i; U) = \text{ReLU}(U_1^T \cdot x)$
2. Two Layers: $f_i(x, y_i; U) = \text{ReLU}(U_2^T \cdot \text{ReLU}(U_1^T \cdot x))$

- **Second order potential:**

1. Linear: $f_{i,j}(x, y_i, y_j; W) = W_{y_i y_j}$



- ▶ **Task:** Image segmentation
- ▶ **Graphical Model:** Fully connected CRF with Gaussian potentials
- ▶ **NN:** PreTrain OxfordNet , predicts 40×40 + upsampling
- ▶ **Inference:** using (algo1), with mean-field as approx. inference








Training method	Mean IoU [%]
Unary only	61.476
Joint	64.060





1. Jointly learning helps
2. Non-linear pairwise function improves over the linear one
3. Deeper and more structured \rightarrow better performance
4. Wide range of applications: Word recognition, Tagging, Segmentation

-  Chen, Schwing, Learning Deep Structured Models [v1](#) [v2](#) [v3](#) [icml](#)
-  Hazan, Schwing, Blending Learning and Infer. in Struct. Pred. [paper](#)
-  Raquel Urtasun, CS Department, UofT, Learning Deep SM [slides](#)
-  Liang-Chieh Chen, CS Department, UofC, ICML [video](#) [slides](#)
-  Alexandr Schwing, CS Department, UofT, Re.Work [video](#)